

CS 591, Lecture 9
Data Analytics: Theory and Applications
Boston University

Babis Tsourakakis

February 22nd, 2017

Announcement

- We will cover the Monday's 2/20 lecture (President's day) by appending half an hour to next Monday's and Wednesday's lectures.
- Next week's office hours for next week: **4.15-5.00pm**


Reminder: Reservoir Sampling

- **Input:** Stream of n items that does not fit in memory $\langle x_1, \dots, x_n \rangle$
- **Goal:** Keep a uniform random sample
- **Simple:** Draw a random integer $\{1, \dots, n\}$, and select that item.
- **Issue:** We don't know the length n of the stream a priori
- **Problem:** Find a uniform sample s from a stream of unknown length

Algorithm: Reservoir Sampling

- 1 Initially we set $s \leftarrow x_1$
- 2 When i -th element arrives, we set $s \leftarrow x_i$ with probability $\frac{1}{i}$

Intuition:



Correctness: What is the probability that $s = x_i$ at some time $t \geq i$?

Algorithm: Reservoir Sampling

Correctness: What is the probability that $s = x_i$ at some time $t \geq i$?

$$\Pr[s = x_i] = \frac{1}{i} \times \left(1 - \frac{1}{i+1}\right) \times \dots \times \left(1 - \frac{1}{t}\right) = \frac{1}{t}.$$

Question: What about the space complexity?

Extension: Weighted Reservoir Sampling


- **Input:** Stream of $\langle w_1, w_2, \dots \rangle$
- **Goal:** Select index i with probability proportional to w_i
- **Again,** we don't know the length n of the stream a priori
- **Define** for each index i

$$W_i = w_1 + \dots + w_i.$$

- **Any ideas?**

Extension: Weighted Reservoir Sampling

Intuition:



Algorithm: Weighted Reservoir Sampling

- 1 Initially we set $s \leftarrow x_1$
- 2 When i -th element arrives, we set $s \leftarrow x_i$ with probability $\frac{w_i}{W_{i-1} + w_i}$

Correctness:

On seeing w_{i+1} :

- We switch to w_{i+1} with probability $\frac{w_{i+1}}{W_{i+1}}$
- If we don't switch the winner $j^* \leq i$ remains with probability

$$\Pr[s = w_{j^*}] = \left(1 - \frac{w_{i+1}}{W_{i+1}}\right) \frac{w_{j^*}}{W_i} = \frac{w_{j^*}}{W_{i+1}}.$$

Question: What about the space complexity?

Reservoir Sampling, k items

- **Input:**

Stream of items $\langle x_1, x_2, \dots, \rangle$


Parameter $k \geq 1$

- **Goal:** Keep k uniform random samples

- **Ideas?**

Extension: Weighted Reservoir Sampling

Intuition:



$$\# \text{ k-sets that do not contain } i = \frac{\binom{i-1}{k}}{\binom{i}{k}} = 1 - \frac{k}{i}.$$

Python implementation

```
import random

def reservoir_sampling(stream,k):
    S = []
    counter = 1
    for x in stream:
        if(counter<=k):
            S.append(x)
        else:
            c = random.randint(0,counter-1);
            if( c<k):
                S[c] = x
            counter +=1
    return S
```

Probabilistic Counting

- Researchers at Bell Labs needed to count overlapping substrings of three letters.
- **Why?** To develop statistic-based spell-checker. This spellchecker (TYPO) was included in early Unix distributions
- Robert Morris came up with an elegant way of counting each item approximately [Morris, 1978]
- **Why approximately?** 8-bit counters can hold counts up to 255 (too small range)...


Probabilistic Counting

- Increments performed in a probabilistic manner
- Suppose we use constant probabilities, i.e., we increase the counter from i to $i + 1$ with probability p
- Why is this problematic?
- Morris' idea: probability $\Pr[i \rightarrow i + 1] = f(i)$ depends on the actual counter value
- Outline: Morris \rightarrow Morris+ \rightarrow Morris++¹.

¹We will follow Jelani Nelson's exposition (M,M+,M++).


Morris' algorithm

- 1 Initialize $X \leftarrow 0$.
- 2 For each event, increment X with probability $\frac{1}{2^X}$.
- 3 Output $\tilde{n} = 2^X - 1$.



Notice that our data structure just maintains an integer! We will see that this integer grows as $\sim \log n$, hence we need $O(\log \log n)$ bits.

Example: Morris' algorithm



Source: [Flajolet, 1985]

$$\begin{aligned}\Pr[\tilde{n} = 1] &= \frac{8}{64}, & \Pr[\tilde{n} = 2] &= \frac{38}{64}, \\ \Pr[\tilde{n} = 3] &= \frac{17}{64}, & \Pr[\tilde{n} = 4] &= \frac{1}{64}.\end{aligned}$$

Morris' algorithm: Analysis

- $X_n :=$ Morris' counter after n updates
- Let's compute $\mathbb{E} [2^{X_n}]$.
- **Idea 1:** Let's apply the formula

$$\mathbb{E} [2^{X_n}] = \sum_{l=0}^n 2^l p_{n,l}.$$

- We have to compute $\mathbf{Pr} [X_n = l] = p_{n,l}$

Morris' algorithm: Analysis

- Suppose we reached state l via n_1 transitions from state 1 to state 1, n_2 from state 2 to state 2, ..., n_l from state l to l

$$(1-2^{-1})^{n_1}2^{-1}(1-2^{-2})^{n_2}2^{-2} \dots (1-2^{-(l-1)})^{n_{l-1}}2^{-(l-1)}(1-2^{-l})^{n_l}.$$

with the condition that $n_1 + \dots + n_l + l - 1 = n$.

- **Tedious approach...**
- Let's try an **inductive** approach.

Morris' algorithm : Analysis

- ① Base case. It's obviously true for $n = 0$.
- ② Induction step.

$$\begin{aligned}\mathbb{E} [2^{X_{n+1}}] &= \sum_{j=0}^{\infty} \mathbf{Pr} [X_n = j] \cdot \mathbb{E} [2^{X_{n+1}} | X_n = j] \\&= \sum_{j=0}^{\infty} \mathbf{Pr} [X_n = j] \cdot (2^j(1 - \frac{1}{2^j}) + \frac{1}{2^j} \cdot 2^{j+1}) \\&= \sum_{j=0}^{\infty} \mathbf{Pr} [X_n = j] \cdot 2^j + \sum_j \mathbf{Pr} [X_n = j] \\&= \mathbb{E} [2^{X_n}] + 1 \\&= (n + 1) + 1\end{aligned}$$

Morris' algorithm : Analysis

- In a similar way, we can show that

$$\mathbb{E} [2^{2X_n}] = \frac{3}{2}n^2 + \frac{3}{2}n + 1.$$

- Therefore

$$\mathbb{V}ar [2^{X_n}] < \frac{1}{2}n^2.$$

- **Chebyshev's inequality:**

$$\mathbf{Pr} [|\tilde{n} - n| \geq \epsilon n] < \frac{1}{2\epsilon^2}.$$

Details

For the second moment, again we use induction:

$$\begin{aligned}\mathbb{E} [2^{2X_n}] &= \sum_{j \geq 0} 2^{2j} \mathbf{Pr} [X_n = j] \\&= \sum_{j \geq 0} 2^{2j} \left(\mathbf{Pr} [X_{n-1} = j] (1 - 2^{-j}) \right. \\&\quad \left. + \mathbf{Pr} [X_{n-1} = j - 1] 2^{-(j-1)} \right) = \dots \\&= \mathbb{E} [2^{2X_{n-1}}] + 3\mathbb{E} [2^{X_{n-1}}] \\&= \frac{3}{2}(n-1)^2 + \frac{3}{2}(n-1) + 1 + 3n \\&= \frac{3}{2}n^2 + \frac{3}{2}n + 1\end{aligned}$$

Morris' algorithm

| | | | | | | | |
|---------------|--------|--------|--------|--------|--------|--------|--------|
| counter value | 7 | 8 | 9 | 10 | 11 | 12 | 13 |
| probability | 0.0011 | 0.0602 | 0.3424 | 0.4218 | 0.1538 | 0.0195 | 0.0001 |

Source: [Flajolet, 1985]

- Exact probability distribution for Morris' counter for $n = 1024$
- Variance is not bad, but we want to do better! **Morris+**

Morris+ algorithm

- We instantiate s independent copies of Morris and average their outputs.
- $s = \frac{1}{2\epsilon^2\delta}$
- Our estimate becomes

$$\tilde{n} = \frac{1}{s} \sum_{i=1}^s \tilde{n}_i.$$

- Now **Chebyshev's inequality** gives

$$\Pr[|\tilde{n} - n| \geq \epsilon n] < \frac{1}{2s\epsilon^2} < \delta.$$

Next step is **Morris++!**

Morris++ algorithm

- Run $t = 36 \lg \frac{1}{\delta}$ copies of **Morris+** with failure probability $\delta = \frac{1}{3}$
- Output the median estimate from the t copies of **Morris+**
- In other words, we have t coin tosses, each results in 1 with probability $\frac{1}{3}$, and in 0 with probability $\frac{2}{3}$.
- Failure if **median** is 0, otherwise **success** → **Chernoff!**

Morris++ algorithm analysis

$$Y_i = \begin{cases} 1, & \text{if } i\text{-th Morris+ fails.} \\ 0, & \text{otherwise.} \end{cases}$$

By Chernoff bound,

$$\Pr \left[\sum_i Y_i > \frac{t}{2} \right] \leq \Pr \left[\sum_i Y_i - \mathbb{E} \left[\sum_i Y_i \right] > \frac{t}{6} \right] < \delta.$$

Morris++ estimate \tilde{n} ($1 \pm \epsilon$)-approximates n with probability at least $(1 - \delta)$.

Question: What is the space complexity that we achieved?

references I



Flajolet, P. (1985).

Approximate counting: a detailed analysis.

BIT Numerical Mathematics, 25(1):113–134.



Morris, R. (1978).

Counting large numbers of events in small registers.

Communications of the ACM, 21(10):840–842.